

SKS Local LLM Mode

A workflow where SKS keeps cloud-grade judgment while a local model on your Mac assists with repeated reading, drafting, and summarization.

- 1** SKS decomposes the work and keeps the safety criteria in place.
- 2** The local LLM handles fast candidate drafts and repeated passes.
- 3** Important decisions and final application still pass through verification.

CONCEPT

The concept is not replacement. It is a helper engine beside SKS.

What the local LLM is good at


- Scanning long files and drafting candidate code or copy
- Repeated summaries, comparison tables, and checklist drafts
- Local pre-exploration and work breakdown support

What SKS still owns

- High-risk decisions around policy, security, and databases
- Current documentation checks and source-backed verification
- Final integration, tests, and Honest Mode reporting

The clearest framing is “fast local drafts plus SKS verification.” Also tell users that final judgment and verification may still process required context through the existing SKS flow.

The workflow is simple



Step 1

Start a local server

Prepare an OpenAI-compatible endpoint with `llama-server`, Ollama, LM Studio, or a similar tool.

Step 2

Connect SKS

Set the local endpoint and model name so SKS can call the helper model.

Step 3

Review before applying

Local drafts are not trusted blindly; they are integrated after tests and review.

Recommended start: run a local server, then enable SKS local mode

1. Prepare the local model server

```
brew install llama.cpp
```

The Hugging Face model cards also document this macOS-friendly flow.

llama-server example

```
llama-server -hf <repo>:<quant>
```

This usually pairs with the `openai-compatible` provider and `http://127.0.0.1:8080`.

2. Connect SKS local mode

```
sks with-local-llm on \  
  --provider openai-compatible \  
  --model <model> \  
  --base-url http://127.0.0.1:8080
```

OpenAI-compatible means SKS calls the local server through an OpenAI-style API.

Check status and turn it off

```
sks with-local-llm status --json  
sks with-local-llm off
```

The setting is stored in `~/.sneakoscope/local-model.json`.

The first run can take time because the model may need to download and populate local caches.

Recommended models by Mac memory

MAC RAM	RECOMMENDATION	MODEL	SHORT NOTE
32GB or more	Primary option	unsloth/Qwen3.6-35B-A3B-MTP-GGUF	User-provided recommendation for Macs with 32GB+ RAM. Use this for users who want the Qwen GGUF and MTP flow.
32GB or more	Alternative option	google/gemma-4-31B-it-qat-q4_0-gguf	User-provided alternative for Macs with 32GB+ RAM. Part of the Google Gemma 4 QAT GGUF family.
Below 32GB	Default starting point	google/gemma-4-12B-it-qat-q4_0-gguf	User-provided default for Macs below 32GB. On 8GB/16GB machines, speed and context length may be limited.

Model launch and SKS connection examples

llama-server: Qwen / Gemma

```
llama-server -hf unsloth/Qwen3.6-35B-A3B-MTP-  
GGUF:UD-Q4_K_M  
  
sks with-local-llm on --provider openai-compatible \  
  --model unsloth/Qwen3.6-35B-A3B-MTP-GGUF:UD-  
Q4_K_M \  
  --base-url http://127.0.0.1:8080
```

Ollama: provider/base URL differences

```
llama-server -hf google/gemma-4-31B-it-qat-q4_0-  
gguf:Q4_0  
  
ollama run hf.co/google/gemma-4-12B-it-qat-q4_0-  
gguf:Q4_0  
  
sks with-local-llm on --provider ollama \  
  --model hf.co/google/gemma-4-12B-it-qat-q4_0-  
gguf:Q4_0 \  
  --base-url http://127.0.0.1:11434
```

Provider and base URL depend on the local server. MLX LM commonly starts with `--provider mlx-lm --base-url http://127.0.0.1:8080`.

Suggested user-facing announcement

Draft announcement copy

SKS now includes a mode for working alongside a local LLM. It keeps SKS verification and orchestration in place while letting a model running on your Mac assist with repeated code reading, summarization, and drafting. For Macs with 32GB+ RAM, start with Qwen3.6 35B or Gemma 4 31B GGUF. Below 32GB, start with Gemma 4 12B GGUF.

Core benefit

Faster repeated drafting and inspection loops

Operating rule

Local output is applied only after verification

Best fit

Developers who use SKS frequently on a Mac

Sources and verified scope

Checked sources

- Hugging Face: `unsloth/Qwen3.6-35B-A3B-MTP-GGUF`
- Hugging Face: `google/gemma-4-31B-it-qat-q4_0-gguf`
- Hugging Face: `google/gemma-4-12B-it-qat-q4_0-gguf`
- getdesign.md: information-first design reference for developer materials

Caveats

- Local model performance depends on the Mac chip, memory, context length, and quantization.
- Sensitive changes, database work, security, and current facts should stay in the SKS verification flow.
- Before starting, run `skt with-local-llm status --json` to check provider, model, and base URL.

Before final distribution, check the SKS local-mode command names and setting keys against the current release docs.